

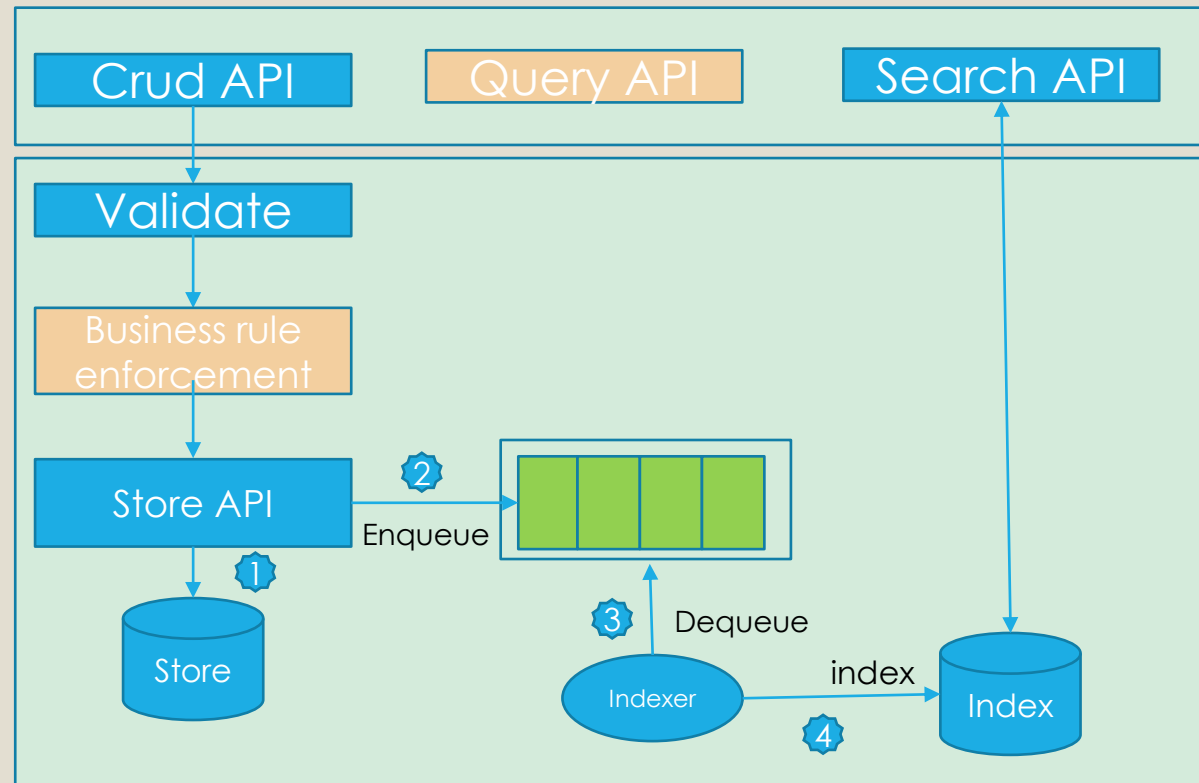


# ADVANCED TOPIC IN BIG DATA

# Quick Review

- By now, you should be familiar with strongly typed data protocols
- You should have reviewed gData, oData, Protocol Buffers
- You should have fair understanding of the overall architecture
- You should have some code working on your laptop

# Architecture



# Prototype Requirements:

Rest API that can handle any structured data in Json

- URIs, status codes, headers, data model, version
- Rest API with support for crd operations
  - Post, Get, Delete
- Rest API with support for validation
  - Json Schema describing the data model for the use case
  - Controller validates incoming payloads against json schema
- The semantics with ReST API operations such as update if not changed/read if changed
  - Update not required
  - Conditional read is required
- Storage of data in key/value store
- Must implement use case provided

# Rest API Specifications


- Data Models
  - Payload structure and serialization
- URI conventions
  - `/ {type} / {id}`
  - `/plan/12xvxc345ssdsds`
- Status Code
  - 200,201
  - 302,304
  - 401, 404, 403, 412, 429
  - 500
- Headers
  - Students should review the HTTP standard headers
  - Various uses of Etag, If-Match, If-None-Match, Authorization in Rest APIs
- Version
  - Accept
  - URL
- Security
- Example: <https://www.hl7.org/fhir/http.html>

# Tooling


- Json simple for Json parsing
- Spring Boot for rest API development
- Elastic Search for search and retrieval capabilities
- Redis for Cache solutions
- Json Schema for schema validation
- Zuul for API Gateway pattern

# But how do I distribute the data?

- single point of failure
- Limited space/storage
- Strongly consistent



- Highly available distributed system
- Seemingly unlimited storage
- What about consistency?



# Key/value stores


- Key readings:
  - Dynamo: Amazon's Highly Available Key-value Store :
    - <http://www.allthingsdistributed.com/files/amazon-dynamo-sosp2007.pdf>
  - Bigtable: A Distributed Storage System for Structured Data:  
<http://static.googleusercontent.com/media/research.google.com/en//archive/bigtable-osdi06.pdf>
  - CAP Theorem
    - Consistency
      - Eventual consistency, Read your own write, Strongly consistent
    - Availability
    - Partition tolerance
    - In the presence of network failure, you have to choose between consistency and high-availability



# Problems

- In the presence of many servers, how do I determine the server that stores the object?
  - Consistent hashing to the rescue
- But what if one of the servers fails or the network connection to the server fails?
  - Replication techniques:
    - Primary/backup
    - Active replication
- If I have multiple servers and if an object is stored on more than one server
  - How do I keep the objects consistent?
    - Eventual consistency, strong consistency, weak consistency

# Weak consistency



# Quorum consistency


R=read replica count

W=write replica count

N=replication factor

Q=**QUORUM** ( $Q = N / 2 + 1$ )

- If  $W + R > N$ , you will have consistency
- On read, two of the replica must respond
- On write, two of the replica must make the data durable before acknowledging the right



# Data Modelling

- $K1 \rightarrow v1$
- $K2 \rightarrow v2$
- $K1?, K2?$

# Consistent hashing

- <http://theory.stanford.edu/~tim/s17/l/l1.pdf>
- How do you map a large number of objects into few servers?
  - $h(x) \bmod n$
- What if the number of servers changes, what would that do to the objects that have already been assigned?
- How do ensure a universal distribution of objects across servers?
- The key idea is:
  - hashing the names of all objects
  - hash the names of all the cache servers  $s$
  - The object and cache names need to be hashed to the same range, such as 32-bit values.

# Key design issues

- Partitioning algorithm
  - Uniform load distribution
- Schema less
- Replication strategy
- Recovering from partial failure
  - Joining a group
  - Load partitioning amongst replicas
- Load rebalancing
- Range query support
- Data versioning
- Support for structured data or simply Blobs
- Marshaling/Unmarshaling
  - How do you store int and floats in redis?

# Mapping of meta-model into key/value store

- JSON payloads can be modeled as a graph.
- [https://www.researchgate.net/publication/315679274\\_Query\\_Service\\_for\\_REST\\_APIs](https://www.researchgate.net/publication/315679274_Query_Service_for_REST_APIs)
- [https://www.researchgate.net/publication/315679444\\_Business\\_Rules\\_for\\_REST\\_APIs](https://www.researchgate.net/publication/315679444_Business_Rules_for_REST_APIs)

How do we map a JSONObject into the key value store?

- What is the key signature?
- Do we store the data as a blob?
- Do we store the data as structured?


# Trade-offs between storing the data as a blob versus structured storage

- Storing data as a blob is fast, atomic, reliable
  - But, how do you update the data?
- Storing the data as structured data requires more work on initial creation, but update are much quicker



# A typical design pattern

A compound document  
with nested objects



Should the compound document be decomposed into its constituent objects for storage, and/or indexing, etc...?